

FUNDAMENTALS OF DATABASES

Simple Queries in SQL

NGUYEN Hoang Ha

Email: nguyen-hoang.ha@usth.edu.vn

Objectives

- Get into the SQL
- Understand the SQL Environment
- Understand the how to write simple queries
- Understand about the string data type, date-time data type
- Understand about NULL values
- Know how to change the data of tables
- Know how to order the output

SQL OVERVIEW

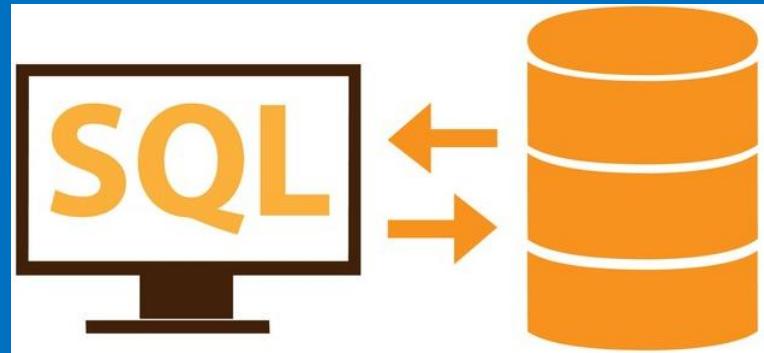
Mathematics to Computer: RA to SQL

Math

π σ ρ \bowtie_c

RA is the conceptual basis for RDB

Computer



What is SQL?

- SEQUEL (Structured English QUERy Language) was developed by IBM in 1974,
later became Structural Query Language (SQL)
- Standard language to work with RDBMS
- Easy to learn
 - Close to English
 - Less than 100 words
- Pronounced as “S-Q-L” or “Sequel.”



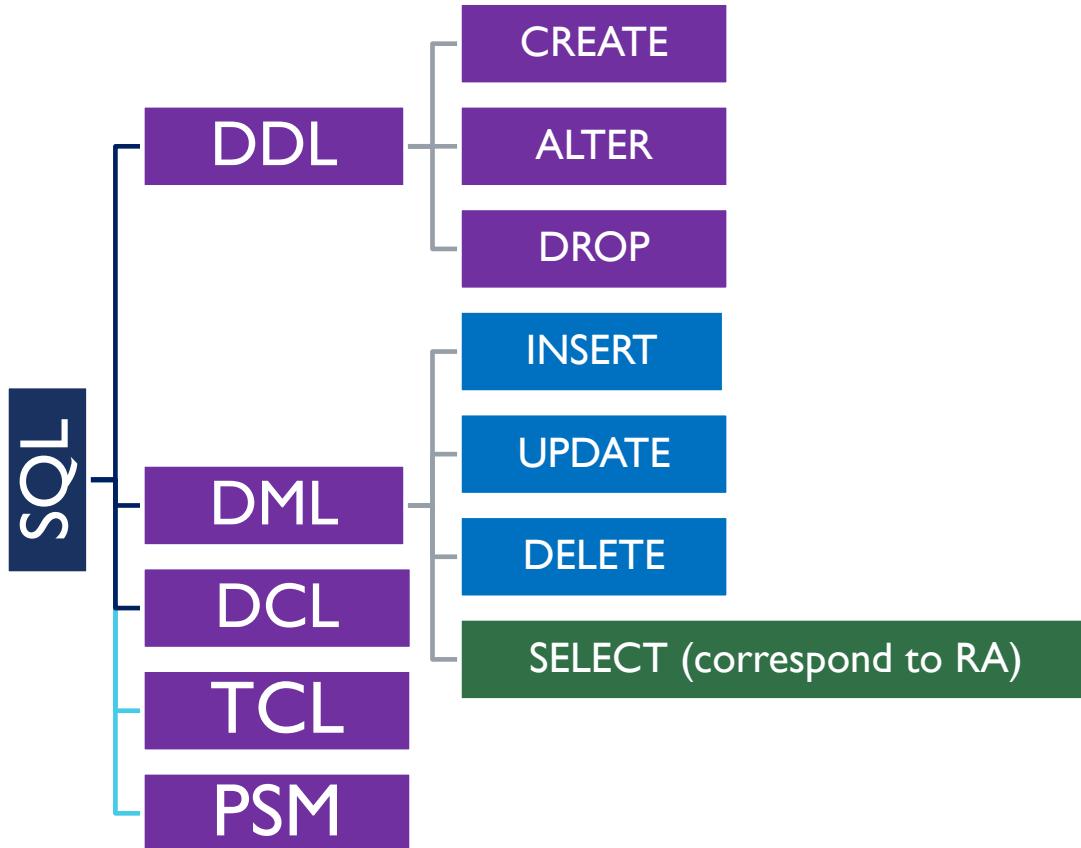
From RA to SQL

- Based on relational algebra, but not entirely identical.
 - Relations \Leftrightarrow Tables
 - Tuples \Leftrightarrow Rows
 - Attributes \Leftrightarrow Columns
- Like a relation, a table is a bag of rows. Duplicates are not automatically removed.
 - This is for practical reasons. Duplicate eliminations are inefficient in implementation.
- Unlike a relation, the order of rows in a table is relevant.

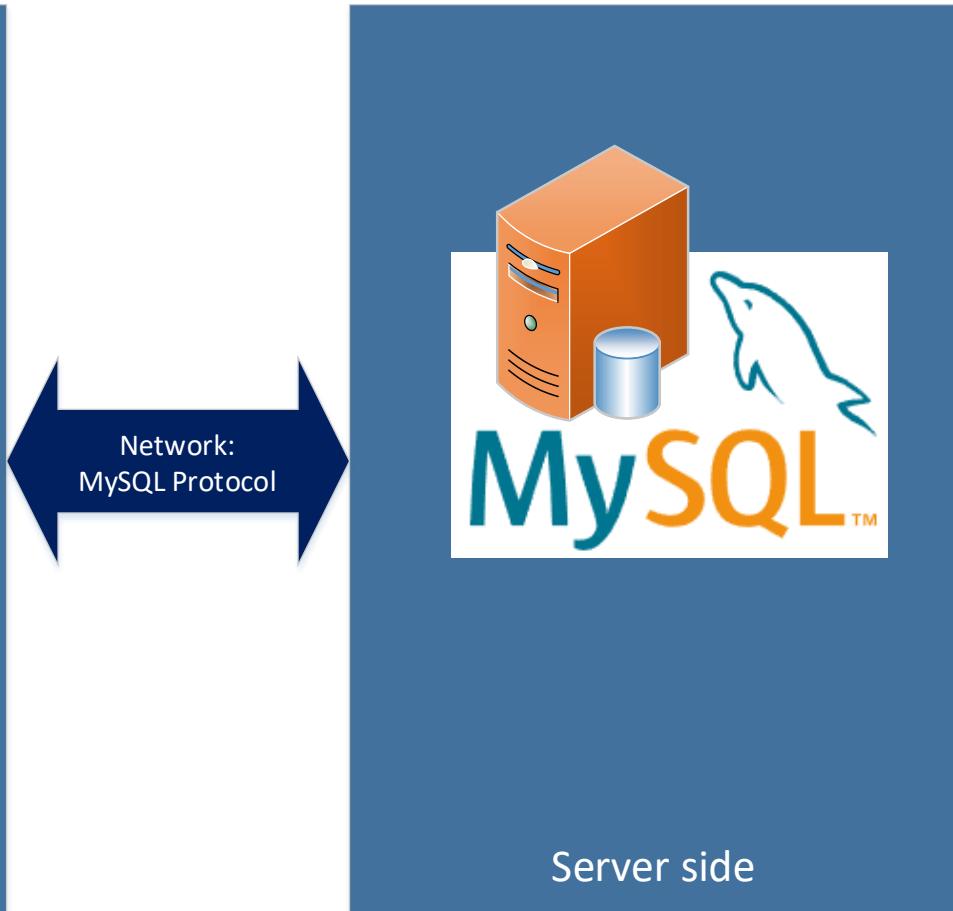
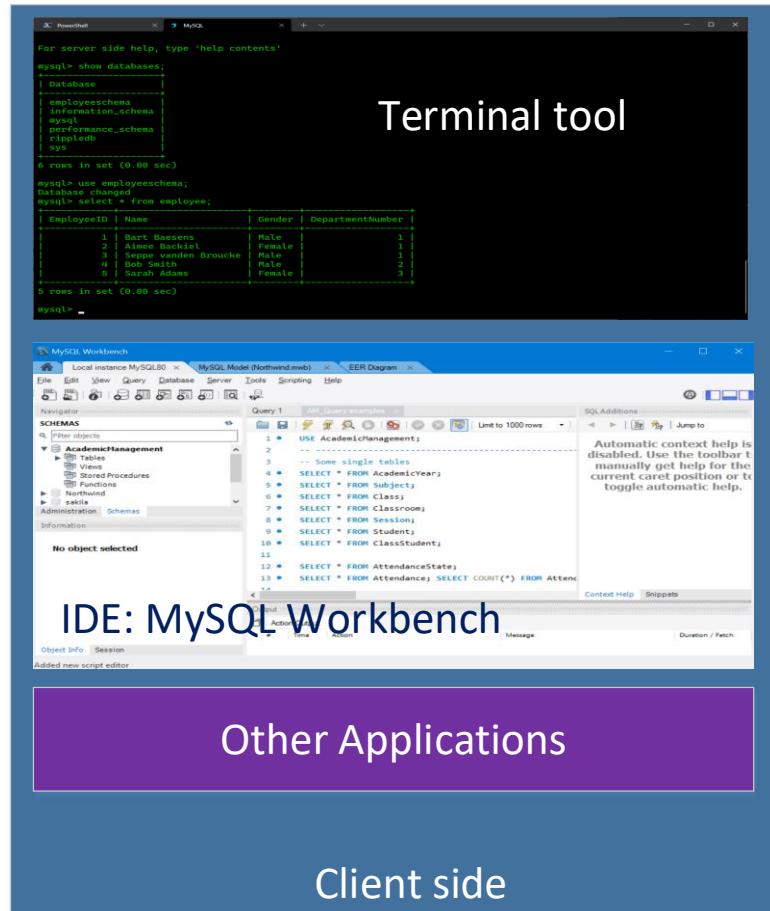
SQL Revisions

| Year | Name | Alias | Comments |
|------|--------------------------|------------------|---|
| 1986 | SQL-86 | SQL-87 | First published by ANSI. Ratified by ISO in 1987. |
| 1989 | SQL-89 | FIPS 127-1 | Minor revision, adopted as FIPS 127-1. |
| 1992 | SQL-92 | SQL2, FIPS 127-2 | Major revision (ISO 9075), <i>Entry Level</i> SQL-92 adopted as FIPS 127-2. |
| 1999 | SQL:1999 | SQL3 | Added regular expression matching, recursive queries, triggers , support for procedural and control-of-flow statements, non-scalar types, and some object-oriented features. |
| 2003 | SQL:2003 | | Introduced XML -related features, <i>window functions</i> , standardized sequences, and columns with auto-generated values (including identity-columns). |
| 2006 | SQL:2006 | | ISO/IEC 9075-14:2006 defines ways in which SQL can be used in conjunction with XML. It defines ways of importing and storing XML data in an SQL database, manipulating it within the database and publishing both XML and conventional SQL-data in XML form. In addition, it provides facilities that permit applications to integrate into their SQL code the use of XQuery , the XML Query Language published by the World Wide Web Consortium (W3C), to concurrently access ordinary SQL-data and XML documents. |
| 2008 | SQL:2008 | | Defines more flexible windowing functions, clarifies SQL 2003 items that were still unclear [1] |

Sub-languages of SQL



SQL Environment



SQL Commands Are Sequential

- Commands are executed in the order they are encountered.
- DDL commands are *not* like C/Java declarations.
- DDL and DML commands can be mixed
 - For example, you can define a table, fill it up with contents, and delete a columns.
 - That is, table definitions (relation schema) can be changed during the lifespan of a database.
 - The ability of doing so does imply it is a good practice.
 - It is best the schema/design of a database is well thought through before its use.

Element of SQL Code

- **Statement:** A complete instruction to the database

```
SELECT * FROM Employees  
WHERE LastName = 'Smith'  
ORDER BY FirstName;
```

- **Clause:** a part of an SQL statement
 - Eg: SELECT * FROM Employees
 - A statement may consist one or many clauses
- **Keyword:** reserved words in SQL having a predefined meaning
- **Identifier:** objects in a databases
 - Eg: tables, columns, indexes, views, schemas, and other elements

Coding in SQL

- SQL is case insensitive
- Convention:
 - Keywords are all in UPPER CASE
 - Identifier name: vary on each project and team, but must be consistent
 - Tables and Columns:
 - PascalCase, e.g.:UsthStudent, InvoiceDetails, TeachingLog)
 - or lowercase with underscore to separate words. e.g.: uesth_student, invoice_details, teaching_log
 - Tables: Singular nouns (customer, order)
 - Columns: Descriptive names (first_name, order_date).
 - Variables and function: use camelCase, e.g.: computeStudentNumber

Coding in SQL (cont’)

- Indentation and Alignment
 - Align SQL statements and clauses for better readability. Each clause normally is written in a line.
 - Use indentation to indicate nested queries

```
SELECT column_name
FROM (
    SELECT column_name
    FROM table_name
    WHERE condition
) AS subquery;
```

Coding in SQL (cont’)

■ Comments

- Use inline comments for complex logic or calculations.

```
SELECT column_name -- This is a comment  
FROM table_name;
```

- Use block comments for detailed explanations.

```
/*  
    This is a block comment.  
    It can span multiple lines.  
*/  
SELECT column_name  
FROM table_name;
```

DDL

DDL Commands

CREATE DATABASE

CREATE TABLE

ALTER TABLE

RENAME TABLE

DROP TABLE

CREATE INDEX

DROP INDEX

Also – CREATE VIEW

EMPLOYEE

| | | | | | | | | | |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|
| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|

DEPARTMENT

| | | | |
|-------|----------------|---------|----------------|
| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|-------|----------------|---------|----------------|

DEPT_LOCATIONS

| | |
|----------------|------------------|
| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|

PROJECT

| | | | |
|-------|----------------|-----------|------|
| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-------|----------------|-----------|------|

WORKS_ON

| | | |
|-------------|------------|-------|
| <u>Essn</u> | <u>Pno</u> | Hours |
|-------------|------------|-------|

DEPENDENT

| | | | | |
|-------------|----------------|-----|-------|--------------|
| <u>Essn</u> | Dependent_name | Sex | Bdate | Relationship |
|-------------|----------------|-----|-------|--------------|

Referential Integrity Constraints

- A referential integrity constraint is used to link (or reference) relations. This means that a foreign key in a relation must also exist in the relation in which it serves as the primary key
 - Super_ssn must be found in Ssn of Employee
 - Mgr_ssn of Department must exist in Ssn of Employee
 - Dno of Employee must exist in Dnumber of Department
 - Dnum of Project must exist in Dnumber of Department
 - Dnumber of Dept_Location must exist in Dnumber of Department

Create Database Example

- To create

```
CREATE DATABASE Company;
```

- To use (or switch to) the database

```
USE Company;
```

- Subsequent commands will operate on the Company database by default.

CREATE TABLE

```
CREATE TABLE base-table-name (colname  
datatype [column constraints - NULL/NOT  
NULL, DEFAULT..., UNIQUE, CHECK..., PRIMARY  
KEY],  
[, colname datatype [column constraints  
...]]]  
.  
.  
.  
[table constraints - PRIMARY KEY..., FOREIGN  
KEY..., UNIQUE..., CHECK...]  
[storage specifications]);
```

Datatypes

- Each column must have a datatype specified
- Standards include various numeric types, fixed-length and varying-length character strings, bit strings, and user-defined types
- Available datatypes vary from DBMS to DBMS

Datatypes

- **char(n).** Fixed length character string, with user-specified length n .
- **varchar(n).** Variable length character strings, with user-specified maximum length n .
- **int.** Integer (a finite subset of the integers that is machine-dependent).
- **smallint.** Small integer (a machine-dependent subset of the integer domain type).
- **numeric(p,d).** Fixed point number, with user-specified precision of p digits, with d digits to the right of decimal point. (ex., **numeric(3,1)**, allows 44.5 to be stored exactly, but not 444.5 or 0.32)
- **real, double precision.** Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(n).** Floating point number, with user-specified precision of at least n digits.
- **Date:** Made up of year-month-day in the format yyyy-mm-dd
- **Time:** Made up of hour:minute:second in the format hh:mm:ss
- **Timestamp:** Has both DATE and TIME components
- **Others:** Boolean, Float, Double Precision
- See user's manual for more data types.

CREATE TABLE Example

```
CREATE TABLE Department (
    Dname          VARCHAR(10) NOT NULL,
    Dnumber        INTEGER    DEFAULT 0,
    Mgr_ssn        CHAR(9),
    Mgr_Start_date CHAR(9),
    PRIMARY KEY   (Dnumber),
    UNIQUE        (Dname),
    FOREIGN KEY   (Mgr_ssn) REFERENCES Employee (Ssn)
);
```

- The “UNIQUE” clause specifies secondary keys.
- Employee has to be created first for the FK Mgr_ssn to refer to it.
- How could we have defined the Dno FK in Employee?

Adding the Dno FK to Employee

- If “CREATE TABLE Employee” is issued first, we cannot specify Dno as a FK in that CREATE command.
- Statement order:
 - CREATE TABLE Employee
 - CREATE TABLE Department
 - Use an ALTER TABLE to add a Foreign key to Employee table

ALTER TABLE Employee

ADD CONSTRAINT

FOREIGN KEY (Dno)

REFERENCES Department (Dnumber) ;

The Check Clause

- Used to specify user-defined constraints
- Assume that dept. numbers are from 0 to 99.

```
CREATE TABLE Department (
    ...
    Dnumber INTEGER Default 0
        CHECK (Dnumber>=0 AND Dnumber<=99) ,
    ...);
```

- “Check” can also be a clause of the entire table.

```
CREATE TABLE Department (
    ...
    Dept_create_date date,
    Mgr_start_date date,
    CHECK (Dept_create_date <= Mgr_start_date)
);
```

Review: Multi attribute Key

- The bar and beer together are the key for Sells:

```
CREATE TABLE Sells (
    bar      CHAR(20) ,
    beer     VARCHAR(20) ,
    price    REAL,
    PRIMARY KEY (bar, beer)
) ;
```

Exercise

- Create the table WORKS_ON, assuming tables EMPLOYEE and PROJECT have been created and Hours ranges from 1 to 56.

Add Columns to Existing Tables

- To add spouse SSN (`S_ssn`) to Employee

```
ALTER TABLE Employee
```

```
ADD COLUMN S_ssn char(9);
```

- The new attribute will have NULLs in all the tuples of the relation right after the command is executed
- Alternatively, we can set a default value.

```
ALTER TABLE EMPLOYEE ADD COLUMN S_ssn  
CHAR(9) DEFAULT "000000000";
```

Delete Columns from Existing Tables

- To delete column S_ssn

```
ALTER TABLE Employee DROP COLUMN S_ssn;
```

- **Reminder:** changing relation schemas typically indicates ill-executed design phase of the database.

Referential Integrity Options

- **Causes** of referential integrity violation for a foreign key FK (consider the Mgr_ssn of Department).
 - **On Delete**: when deleting the foreign tuple
 - What to do when deleting the manager tuple in Employee ?
 - **On Update**: when updating the foreign tuple
 - What to do when updating/changing the SSN of the manager tuple in Employee is changed ?
- **Actions** when the above two causes occur.
 - **Set Null**: the Mgr_ssn is set to null.
 - **Set Default**: the Mgr_ssn is set to the default value.
 - **Cascade**: the Mgr_ssn is updated accordingly
 - If the manager is deleted, the department is also deleted.

The Mgr_ssn Example

```
CREATE TABLE DEPARTMENT (
    ...
    Mgr_ssn      CHAR(9),
    ...
    FOREIGN KEY (Mgr_ssn)
        REFERENCES EMPLOYEE (Ssn)
        ON DELETE    ???
        ON UPDATE    ???
) ;
```

Another Example

```
CREATE TABLE EMP (
    ...
    SSN      CHAR(9),
    DNO      INTEGER DEFAULT 1,
    SUPERSSN CHAR(9),
    PRIMARY KEY (ESSN),
    FOREIGN KEY (DNO) REFERENCES DEPT
        ON DELETE SET DEFAULT
        ON UPDATE CASCADE,
    FOREIGN KEY (SUPERSSN) REFERENCES EMP
        ON DELETE SET NULL
        ON UPDATE CASCADE
);
```

Miscellaneous Commands

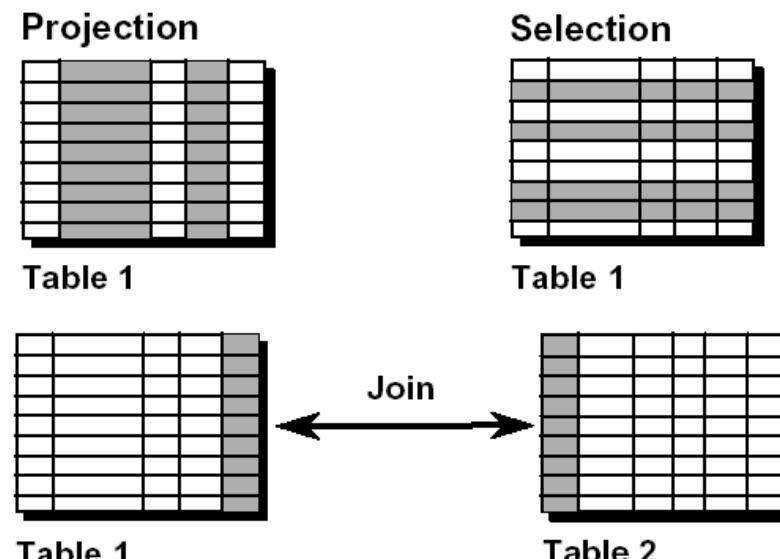
- SHOW DATABASES ;
 - Show all the databases on the server
- SHOW TABLES ;
 - Show all the tables of the present database
- DROP TABLE *t_name* ;
 - Delete the entire table *t_name*
- DROP DATABASE *db_name* ;
 - Delete the entire database *db_name*
- DESCRIBE *table_name*
 - Show structure of table, only for MySQL

SIMPLE DML QUERIES

SELECT commands

A SELECT statement retrieves information from the database. Using a SELECT statement, you can do the following:

- **Projection:** You can use the projection capability in SQL to choose the columns in a table that you want returned by your query.
- **Selection:** You can use the selection capability in SQL to choose the rows in a table that you want returned by a query (with WHERE clause)
- **Joining:** You can use the join capability in SQL to bring together data that is stored in different tables by creating a link between them.



SQL data retrieval query structure

| | | | |
|------------------|-----------------------------------|-------------|-------------|
| SELECT | desired expressions, columns | π | ρ |
| [FROM | one or more tables] | \bowtie_C | |
| [WHERE | Conditions about expected rows] | σ | |
| [GROUP BY | rows with the same column values] | | Next lesson |
| [ORDER BY | column list] | | |

Syntax for a simple SELECT queries

- ❖ **SELECT** clause identifies *what* columns
 - **ALL**: Specifies that duplicate rows can appear in the result set. ALL is the default
 - **DISTINCT**: Specifies that only unique rows can appear in the result set. Null values are considered equal for the purposes of the DISTINCT keyword
 - **TOP *n* [PERCENT]**:Specifies that only the first *n* rows are to be output from the query result set. *n* is an integer between 0 and 4294967295. If PERCENT is also specified, only the first *n* percent of the rows are output from the result set. When specified with PERCENT, *n* must be an integer between 0 and 100
- ❖ **FROM** identifies *which* table

A trick for reading & writing queries

- It's generally easiest to examine a **SELECT-FROM-WHERE** query by:
 - First looking at the **FROM** clause to learn which relations are involved in the query
 - Then, move to the **WHERE** clause to learn what it is about tuples that is important to the query
 - Finally, look at the **SELECT** clause to see what the output format is
- The same order: **FROM**, then **WHERE**, then **SELECT** is often useful when writing queries of your own as well

Example: SELECT all columns

The "Persons" table:

| P_Id | LastName | FirstName | Address | City |
|------|-----------|-----------|--------------|-----------|
| 1 | Hansen | Ola | Timoteivn 10 | Sandnes |
| 2 | Svendson | Tove | Borgvn 23 | Sandnes |
| 3 | Pettersen | Kari | Storgt 20 | Stavanger |

Now we want to select all the columns from the "Persons" table.

We use the following SELECT statement:

```
SELECT * FROM Persons
```

Tip: The asterisk (*) is a quick way of selecting all columns!

The result-set will look like this:

| P_Id | LastName | FirstName | Address | City |
|------|-----------|-----------|--------------|-----------|
| 1 | Hansen | Ola | Timoteivn 10 | Sandnes |
| 2 | Svendson | Tove | Borgvn 23 | Sandnes |
| 3 | Pettersen | Kari | Storgt 20 | Stavanger |

Projection in SQL

- We can, if we wish, eliminate some of the components of the chosen tuples; that is, we can project the relation produced by a SQL query onto some of its attributes
- In place of the * of the SELECT clause, we may list some of the attributes of the relation mentioned in the FROM clause. The result will be projected onto the attributes listed

Example: Projection in SQL

The "Persons" table:

| P_Id | LastName | FirstName | Address | City |
|------|-----------|-----------|--------------|-----------|
| 1 | Hansen | Ola | Timoteivn 10 | Sandnes |
| 2 | Svendson | Tove | Borgvn 23 | Sandnes |
| 3 | Pettersen | Kari | Storgt 20 | Stavanger |

Now we want to select the content of the columns named "LastName" and "FirstName" from the table above. We use the following SELECT statement:

```
SELECT LastName, FirstName FROM Persons
```

The result-set will look like this:

| LastName | FirstName |
|-----------|-----------|
| Hansen | Ola |
| Svendson | Tove |
| Pettersen | Kari |

Example: Extended projection using Arithmetic Operators

```
SELECT last_name, salary, salary + 300
FROM   employees;
```

| LAST_NAME | SALARY | SALARY+300 |
|-----------|--------|------------|
| King | 24000 | 24300 |
| Kochhar | 17000 | 17300 |
| De Haan | 17000 | 17300 |
| Hunold | 9000 | 9300 |
| Ernst | 6000 | 6300 |
| Lorentz | 4200 | 4500 |

- Note that the resultant calculated column SALARY+300 is not a new column in the EMPLOYEES table; it is for display only.
- By default, the name of a new column comes from the calculation that generated it—in this case, salary+300.

Renaming or Defining a Column Alias

- **A column alias:**

- Renames a column heading
- Is useful with calculations
- Immediately follows the column name - there can also be the optional AS keyword between the column name and alias

Example: ALIAS

- The example displays the last names and annual salaries of all the employees.
- Because Annual Salary contain a space, it has been enclosed in double quotation marks.
- Notice that the column heading in the output is exactly the same as the column alias.

```
SELECT last_name "Name",
       salary*12 "Annual Salary"
FROM   employees;
```

| Name | Annual Salary |
|---------|---------------|
| King | 288000 |
| Kochhar | 204000 |
| Higgins | 144000 |
| Gietz | 96600 |

Duplication Eliminating with SELECT distinct

The "Persons" table:

| P_Id | LastName | FirstName | Address | City |
|------|-----------|-----------|--------------|-----------|
| 1 | Hansen | Ola | Timoteivn 10 | Sandnes |
| 2 | Swendson | Tove | Borgvn 23 | Sandnes |
| 3 | Pettersen | Kari | Storgt 20 | Stavanger |

Now we want to select only the distinct values from the column named "City" from the table above. We use the following SELECT statement:

```
SELECT DISTINCT City FROM Persons
```

The result-set will look like this:

| City |
|-----------|
| Sandnes |
| Stavanger |

Selection in SQL or Restricting data

- While retrieving data from the database, you may need to restrict the rows of data that are displayed
- In that case, the solution is to use the WHERE clause
- The WHERE clause is equal to the selection operator of relational algebra
- The expression that may follow WHERE include conditional expressions like those found in C or Java

Selection in SQL (or Restricting data)

```
SELECT [ ALL | DISTINCT ]
    [ TOP n[ PERCENT ] ]
    * | {column_name | expression [alias],...}
[FROM table]
[WHERE conditions]
```

- WHERE: restricts the query to rows that meet a condition
- The WHERE clause follows the FROM clause.
- *Condition*: is composed of column names, expressions, constants, and a comparison operator

Example: Restricting data

```
SELECT employee_id, last_name, job_id, department_id  
FROM   employees  
WHERE  department_id = 90;
```

```
SELECT last_name, job_id, department_id  
FROM   employees  
WHERE  last_name = 'Goyal';
```

```
SELECT last_name, salary  
FROM   employees  
WHERE  salary <= 3000;
```

```
SELECT last_name, salary  
FROM   employees  
WHERE  salary BETWEEN 2500 AND 3500;
```

Example: Restricting data

```
SELECT employee_id, last_name, salary, manager_id  
FROM   employees  
WHERE  manager_id IN (100, 101, 201);
```

```
SELECT first_name  
FROM   employees  
WHERE  first_name LIKE 'S%';
```

```
SELECT last_name, manager_id  
FROM   employees  
WHERE  manager_id IS NULL;
```

```
SELECT employee_id, last_name, job_id, salary  
FROM   employees  
WHERE  salary >=10000  
AND    job_id LIKE '%MAN%';
```

Example: Restricting data

```
SELECT last_name, job_id  
FROM   employees  
WHERE  job_id NOT IN ('IT_PROG', 'ST_CLERK', 'SA_REP');
```

Comparison of Strings

- Two strings are equal if they are the same sequence of characters. Recall from the section 2.3.2 that strings can be stored as fixed-length strings (using CHAR) or variable-length strings (using VCHAR)
- When comparing strings with different declarations, only the actual strings are compared (SQL ignores any “pad” characters that must be present in the database in order to give a string its required length)
- We can use “<”, “>”, “=” and “<>” operators to compare two strings

Pattern matching in SQL

- SQL also provides the capability to compare strings on the basis of a simple pattern match. An alternative form of comparison expression is:

s LIKE p

where:

- S: is a string
- P: is a pattern (with the optional use of some special characters: "%", "_" ..)
- Similarly, "s NOT LIKE p" is true if and only if string s does not match pattern p

Wildcards

SQL Server

MySQL

| Symbol | Meanings | Example |
|--------|---------------------------------------|---------------------------------|
| % | Any string of zero or more characters | “P%” finds Peter, pack |
| _ | Underscore for any single character | “h_t” finds hot, hat, hit |
| [] | Any single character within brackets | “h[oa]t” finds hot, hat |
| [^] | Any single character not in brackets | “h[^oa]t” finds hit but not hot |
| - | A range of characters | “c[a-f]%” finds cat, cell |

Dates and Times

- SQL generally support dates and times as special data types. These values are often representable in a variety of formats such as:
 - '05/14/1948' or
 - '14 May 1948'
- We can compare dates or times using the same comparison operators we use for numbers or strings

NULL values

- **Null** means 'nothing' or without value or consequence
- **Null** is a special marker used in Structured Query Language (SQL) to indicate that a data value does not exist in the database. Introduced by the creator of the relational database model
- Since Null is not a member of any data domain, it is not considered a "value", but rather a marker (or placeholder) indicating the absence of value. Because of this, comparisons with Null can never result in either True or False, but always in a third logical result, Unknown
- However, certain operations on Null can return values if the value of Null is not relevant to the outcome of the operation

Ordering the Output

- While retrieving data from the database, you may need to specify the order in which the rows are displayed.
- In that case, the solution is to use the ORDER BY clause

Ordering the Output

- If you use the ORDER BY clause, it must be the last clause of the SQL statement.
- **Expression:** Specifies a column on which to sort.
A sort column can be specified as a name or column alias (which can be qualified by the table or view name), an expression, or a nonnegative integer representing the position of the name, alias, or expression in select list.
- Multiple sort columns can be specified. The sequence of the sort columns in the ORDER BY clause defines the organization of the sorted result set.

```
SELECT [ ALL | DISTINCT ]  
      [ TOP n[ PERCENT ] ]  
      * | {column_name | expression [alias],...}  
FROM table  
[WHERE conditions]  
[ORDER BY {expression [ASC | DESC] ,...} ]
```

Example: Sorting data

```
SELECT last_name, job_id, department_id, hire_date  
FROM employees  
ORDER BY hire_date DESC;
```

Sorting by Column Alias

```
SELECT employee_id, last_name, salary*12 annsal  
FROM employees  
ORDER BY annsal;
```

Sorting by Multiple Columns

```
SELECT last_name, department_id, salary  
FROM employees  
ORDER BY department_id, salary DESC;
```

Exercise I

Write SQL queries to create the following tables:

- Studio (name, address)
- Star (name, address, phone)
- Movie (title, year, length, genre)

After creating, write SQL queries to drop them

Exercise 2

Write SQL queries to do following tasks:

- Add a column named Description into Movie table
(you must determine the data type for it)

- Add a column named Hobbies into Star table (you
must determine the data type for it)

- Add a column named Birthdate into Star table (you
must determine the data type for it)

Exercise 3

Write SQL queries to do following tasks:

- Remove the column named Description from Movie table
- Remove the column named Hobbies from Star table
- Remove the column named Birthdate from Star table

Exercise 4

- Write a SQL query to show all tuples in table Employees

EMPLOYEES

| LAST_NAME | DEPARTMENT_ID | SALARY |
|-----------|---------------|--------|
| Getz | 10 | 3000 |
| Davis | 20 | 1500 |
| King | 20 | 2200 |
| Davis | 30 | 5000 |
| Kochhar | | 5000 |

Exercise 5

- Write a SQL query to show all SALARY (but eliminating duplicates) in table EMPLOYEES
- Write a SQL query to show all DEPARTMENT_ID (but eliminating duplicates) in table EMPLOYEES

EMPLOYEES

| LAST_NAME | DEPARTMENT_ID | SALARY |
|-----------|---------------|--------|
| Getz | 10 | 3000 |
| Davis | 20 | 1500 |
| King | 20 | 2200 |
| Davis | 30 | 5000 |
| Kochhar | | 5000 |

DML STATEMENTS

UPDATE INSERT DELETE

Reference: Section 6.5 Jeffrey D. Ullman, Jennifer Widom: A First Course in Database Systems, Pearson, 3rd Edition (2007)

Insert record

```
INSERT INTO table_name ( field1, field2,...fieldN )  
VALUES  (value1, value2,...valueN );
```

```
INSERT INTO table_name ( field1, field2,...fieldN )  
VALUES  (value1, value2,...valueN ),  
        ( value1, value2,...valueN );
```

Insert data from a select query

```
INSERT INTO table_name ( field1, field2,...fieldN )  
<SelectStatement>
```

```
INSERT INTO SeattleEmp(id, last_name, first_name)  
SELECT id, last_name, first_name  
FROM employees  
WHERE city = 'Seattle'
```

Update records

UPDATE TableName

SET column1 = value1, column2 = value2,.

WHERE condition

UPDATE employees

SET city = 'SEAT'

WHERE city = 'seattle'

Delete records

```
DELETE FROM TableName  
[WHERE condition]
```

```
DELETE FROM student WHERE name='Shannon';
```

```
DELETE FROM student; -- delete all students
```

Exercise 6

- Write a SQL query to delete all tuples in EMPLOYEES table
- Write a SQL query to delete all tuples with NULL value in DEPARTMENT_ID

EMPLOYEES

| LAST_NAME | DEPARTMENT_ID | SALARY |
|-----------|---------------|--------|
| Getz | 10 | 3000 |
| Davis | 20 | 1500 |
| King | 20 | 2200 |
| Davis | 30 | 5000 |
| Kochhar | | 5000 |

Exercise 7

- Write a SQL query to set **DEPARTMENT_ID** to the value 10
- Write a SQL query to set **DEPARTMENT_ID** to the value 10 if **DEPARTMENT_ID** is **NULL**

EMPLOYEES

| LAST_NAME | DEPARTMENT_ID | SALARY |
|-----------|---------------|--------|
| Getz | 10 | 3000 |
| Davis | 20 | 1500 |
| King | 20 | 2200 |
| Davis | 30 | 5000 |
| Kochhar | | 5000 |

Exercise 8

- Write a SQL query to insert some new tuples into EMPLOYEES

EMPLOYEES

| LAST NAME | DEPARTMENT ID | SALARY |
|-----------|---------------|--------|
| Getz | 10 | 3000 |
| Davis | 20 | 1500 |
| King | 20 | 2200 |
| Davis | 30 | 5000 |
| Kochhar | | 5000 |